

Specification and Implementation of Replicated List

— The Jupiter Protocol Revisited

(Brief Announcement at PODC'2018)

Hengfeng Wei, Yu Huang, Jian Lu

Nanjing University

July 24, 2018



Brief Announcement

The **Jupiter protocol** [Nichols et al., 1995]^a for replicated list **satisfies** the **weak list specification** [Attiya et al., 2016]^b.

^aDavid A. Nichols et al. (1995). “High-latency, Low-bandwidth Windowing in the Jupiter Collaboration System”. In: *Proceedings of the 8th Annual ACM Symposium on User Interface and Software Technology*. UIST '95. ACM, pp. 111–120.

^bHagit Attiya et al. (2016). “Specification and complexity of collaborative text editing”. In: *Proceedings of the 2016 ACM Symposium on Principles of Distributed Computing*. PODC '16. ACM, pp. 259–268.

This was proposed as a conjecture in a PODC paper [Attiya et al., 2016].

Replicated List

Replicated, highly available collaborative text editing systems



(a) Google Docs



(b) Apache Wave



(c) Wikipedia



(d) L^AT_EX Editor

Replicas are required to respond to user operations **immediately**.
Updates are propagated to other replicas **asynchronously**.

Replicated list object: to model the core functionality

$\text{INS}(a, p)$: Insert a at position p .

$\text{DEL}(p)$: Delete the element at position p .

READ : Return the list.

Weak List Specification

Specification and Complexity of Collaborative Text Editing

Hagit Attiya
Technion

Sebastian Burckhardt
Microsoft Research

Alexey Gotsman
IMDEA Software Institute

Adam Morrison
Technion

Hongseok Yang
University of Oxford

Marek Zawirski*
Inria & Sorbonne Universités,
UPMC Univ Paris 06, LIP6




Definition (Weak List Specification $\mathcal{A}_{\text{weak}}$ [Attiya et al., 2016])

Informally, $\mathcal{A}_{\text{weak}}$ requires the ordering between **elements that are not deleted** to be consistent across the system.

Specify a global property on all states across the system.


Pairwise state compatibility property:

For any *pair* of list states, there *cannot* be two elements a and b such that a precedes b in one state but b precedes a in the other.



Jupiter


Jupiter adopts the **client-server** architecture [Nichols et al., 1995]:



Operations are **totally ordered** at the server.

Client $\xrightarrow{\text{FIFO}}$ server $\xrightarrow{\text{FIFO}}$ other clients


To achieve convergence, Jupiter uses $2D$ state spaces [Xu, Sun, and Li, 2014] to manage how and when to perform OTs¹ [Ellis and Gibbs, 1989].



There can be ≤ 2 edges coming from the same node (LOCAL or GLOBAL).

¹OT: Operational Transformation

Each **client** maintains a $2D$ state space.



The **server** maintains $n (= 3)$ $2D$ state spaces, one for each client.


Global property on all replica states specified by $\mathcal{A}_{\text{weak}}$



Local view each replica maintains in Jupiter

CJupiter (Compact Jupiter)

CJupiter maintains an n -ary ordered state space for each replica.




There can be **more than two edges** coming from the same node.

Edges from the same node are **totally ordered** by associated operations.

Proposition (Compactness of CJupiter (Informal))

At a high level, CJupiter maintains only **one** n -ary ordered state space.



Each replica behavior corresponds to a **path** going through this state space.


Theorem (Equivalence of CJupiter and Jupiter)

Under the same schedule, the behaviors of corresponding replicas in CJupiter and Jupiter are the same.

From the perspectives of both the server and the clients.

CJupiter Satisfies the Weak List Specification

We focus on a single n -ary ordered state space.




We show the **pairwise state compatibility** property in **three steps**.

- 1 Take any two nodes/states v_1 and v_2 .

Lemma (LCA (Lowest Common Ancestor))

*Each pair of states in the n -ary ordered state space has a **unique** LCA.*




$$v_0 = \text{LCA}(v_1, v_2)$$

- 2 Consider the paths to v_1 and v_2 from their LCA v_0 .

Lemma (Disjoint Paths)

The set of operations $O_{v_0 \rightsquigarrow v_1}$ along $P_{v_0 \rightsquigarrow v_1}$ is *disjoint* from the set of operations $O_{v_0 \rightsquigarrow v_2}$ along $P_{v_0 \rightsquigarrow v_2}$.




$$v_0 = \text{LCA}(v_1, v_2)$$

- 3 Consider the states in these two paths.

Lemma (Compatible Paths)

Each pair of states consisting of one state v in $P_{v_0 \rightsquigarrow v_1}$ and the other v' in $P_{v_0 \rightsquigarrow v_2}$ are **compatible**.



$$v_0 = \text{LCA}(v_1, v_2)$$

In particular,
 v_1 and v_2 are compatible.

Brief Announcement

The **Jupiter** protocol [Nichols et al., 1995]^a for replicated list **satisfies** the **weak list specification** [Attiya et al., 2016]^b.

^aDavid A. Nichols et al. (1995). “High-latency, Low-bandwidth Windowing in the Jupiter Collaboration System”. In: *Proceedings of the 8th Annual ACM Symposium on User Interface and Software Technology*. UIST '95. ACM, pp. 111–120.

^bHagit Attiya et al. (2016). “Specification and complexity of collaborative text editing”. In: *Proceedings of the 2016 ACM Symposium on Principles of Distributed Computing*. PODC '16. ACM, pp. 259–268.

Thank
You!



- Attiya, Hagit et al. (2016). "Specification and complexity of collaborative text editing". In: *Proceedings of the 2016 ACM Symposium on Principles of Distributed Computing*. PODC '16. ACM, pp. 259–268.
- Ellis, C. A. and S. J. Gibbs (1989). "Concurrency Control in Groupware Systems". In: *Proceedings of the 1989 ACM SIGMOD International Conference on Management of Data*. SIGMOD '89. ACM, pp. 399–407.
- Nichols, David A. et al. (1995). "High-latency, Low-bandwidth Windowing in the Jupiter Collaboration System". In: *Proceedings of the 8th Annual ACM Symposium on User Interface and Software Technology*. UIST '95. ACM, pp. 111–120.
- Roh, Hyun-Gul et al. (2011). "Replicated Abstract Data Types: Building Blocks for Collaborative Applications". In: *J. Parallel Distrib. Comput.* 71.3, pp. 354–368.
- Shapiro, Marc et al. (2011). "Conflict-free Replicated Data Types". In: *Proceedings of the 13th International Conference on Stabilization, Safety, and Security of Distributed Systems*. SSS'11. Springer-Verlag, pp. 386–400.
- Xu, Yi, Chengzheng Sun, and Mo Li (2014). "Achieving Convergence in Operational Transformation: Conditions, Mechanisms and Systems". In: *Proceedings of the 17th ACM Conference on Computer Supported Cooperative Work*. CSCW '14. ACM, pp. 505–518.

Backup

Replication (for availability)



Replicas respond to user operations **immediately**

Updates are propagated **asynchronously**

Definition (Eventual Convergence (EC) [Ellis and Gibbs, 1989])

The lists at all replicas are identical *at quiescence*.

Definition (Strong Eventual Consistency (SEC) [Shapiro et al., 2011])

The lists at the replicas that *have executed the same set of user operations* are identical.

Specify little on *intermediate states* going through by replicas.

Strong/weak list specification [Attiya et al., 2016]

Specify global properties on all states across the system.

Specification and Complexity of Collaborative Text Editing

Hagit Attiya
Technion

Sebastian Burckhardt
Microsoft Research

Alexey Gotsman
IMDEA Software Institute

Adam Morrison
Technion


Hongseok Yang
University of Oxford

Marek Zawirski^{*}
Inria & Sorbonne Universités,
UPMC Univ Paris 06, LIP6

Proved: RGA [Roh et al., 2011] satisfies the strong list specification.

Conjecture: Jupiter [Nichols et al., 1995] satisfies the weak list specification.

It is still challenging to achieve convergence despite the server.



Serializability may not be desirable.


It does not imply that clients process operations in the same order.

$$\forall \sigma, \sigma' : a, b \in \sigma \cap \sigma' \implies (a \prec_{\sigma} b \iff a \prec_{\sigma'} b)$$

($\sigma, \sigma' : \text{list}; \quad a, b : \text{element}; \quad \prec_{\sigma} : \text{precedes}$)

$\sigma_0 : ab$


$\sigma_1 : ba$



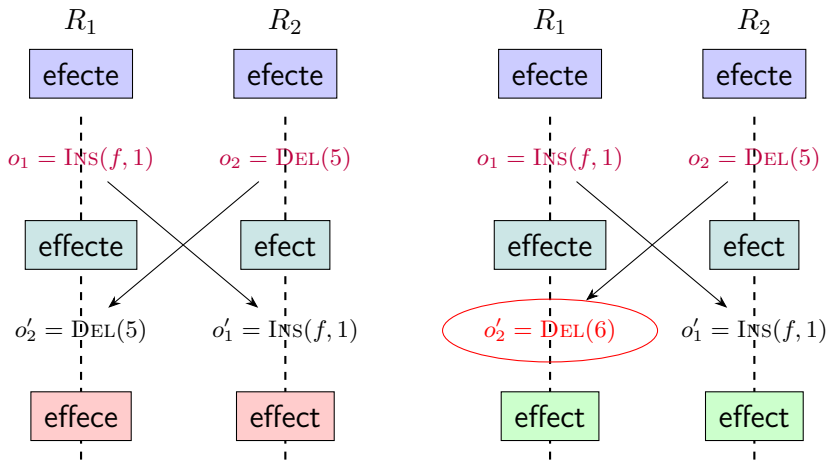
$\sigma_0 : ax$

$\sigma_1 : xb$



$\sigma_2 : ba$



OT (Operational Transformation) [Ellis and Gibbs, 1989]



Jupiter utilizes OT² [Ellis and Gibbs, 1989] to achieve convergence.



$$\sigma; o_1; o_2' \equiv \sigma; o_2; o_1'$$

OT functions for a replicated list object [Ellis and Gibbs, 1989]


$$OT\left(\text{INS}(a_1, p_1, pr_1), \text{INS}(a_2, p_2, pr_2)\right) = \begin{cases} \text{INS}(a_1, p_1, pr_1) & p_1 < p_2 \\ \text{INS}(a_1, p_1 + 1, pr_1) & p_1 > p_2 \\ \text{NOP} & p_1 = p_2 \wedge a_1 = a_2 \\ \text{INS}(a_1, p_1 + 1, pr_1) & p_1 = p_2 \wedge a_1 \neq a_2 \wedge pr_1 > pr_2 \\ \text{INS}(a_1, p_1, pr_1) & p_1 = p_2 \wedge a_1 \neq a_2 \wedge pr_1 \leq pr_2 \end{cases}$$

$$OT\left(\text{INS}(a_1, p_1, pr_1), \text{DEL}(-, p_2, pr_2)\right) = \begin{cases} \text{INS}(a_1, p_1, pr_1) & p_1 \leq p_2 \\ \text{INS}(a_1, p_1 - 1, pr_1) & p_1 > p_2 \end{cases}$$

$$OT\left(\text{DEL}(-, p_1, pr_1), \text{INS}(a_2, p_2, pr_2)\right) = \begin{cases} \text{DEL}(-, p_1, pr_1) & p_1 < p_2 \\ \text{DEL}(-, p_1 + 1, pr_1) & p_1 \geq p_2 \end{cases}$$

$$OT\left(\text{DEL}(-, p_1, pr_1), \text{DEL}(-, p_2, pr_2)\right) = \begin{cases} \text{DEL}(-, p_1, pr_1) & p_1 < p_2 \\ \text{DEL}(-, p_1 - 1, pr_1) & p_1 > p_2 \\ \text{NOP} & p_1 = p_2 \end{cases}$$

Consider a replicated system with n ($= 3$) clients.



Theorem (Equivalence of CJupiter and Jupiter)

Under the same schedule, the behaviors of corresponding replicas in CJupiter and Jupiter are the same.

At the server side:

Proposition ($n \leftrightarrow 1$ (Informal))

*The single n -ary ordered state space at the server side in CJupiter is a **compact representation** of n 2D state spaces at the server side in Jupiter.*

At the client side:

Proposition ($1 \leftrightarrow 1$ (Informal))

*Jupiter is **slightly optimized in implementation** at clients by eliminating redundant OTs than CJupiter.*